Sharon Laurance

Exercise 1: Preparing your Hadoop infrastructure

Exercise 1.1: Setting up a Hadoop infrastructure

Installation of Hadoop in Pseudo-Distributed mode

2. Initialize the HDFS file system by the command: hadoop namenode -format.

```
C:\WINDOWS\system32>hdfs namenode -format
2022-06-01 13:55:33,796 INFO namenode.NameNode: STARTUP_MSG:
/************************************************************
STARTUP_MSG: Starting NameNode
STARTUP_MSG:   host = DESKTOP-9QR3L69/147.172.216.203
STARTUP_MSG:   args = [-format]
STARTUP_MSG:   version = 3.2.3
STARTUP_MSG:   classpath = C:\hadoop-3.2.3\etc\hadoop;C:\hadoop-3.2.3\share\hadoop\common
2.3\share\hadoop\common\lib\animal-sniffer-annotations-1.17.jar;C:\hadoop-3.2.3\share\had
nnotations-0.5.0.jar;C:\hadoop-3.2.3\share\hadoop\common\lib\avro-1.7.7.jar;C:\hadoop-3.2
common\lib\commons-beanutils-1.9.4.jar;C:\hadoop-3.2.3\share\hadoop\common\lib\commons-cl
oop-3.2.3\share\hadoop\common\lib\commons-collections-3.2.2.jar;C:\hadoop-3.2.3\share\had
b\commons-configuration2-2.1.1.jar;C:\hadoop-3.2.3\share\hadoop\common\lib\commons-io-2.8
3.2.3\share\hadoop\common\lib\commons-logging-1.1.3.jar;C:\hadoop-3.2.3\share\hadoop\comm
net-3.6.jar;C:\hadoop-3.2.3\share\hadoop\common\lib\commons-text-1.4.jar;C:\hadoop-3.2.3\
common\lib\curator-framework-2.13.0.jar;C:\hadoop-3.2.3\share\hadoop\common\lib\curator-r
hadoop-3.2.3\share\hadoop\common\lib\error_prone_annotations-2.2.0.jar;C:\hadoop-3.2.3\sh
\lib\gson-2.2.4.jar;C:\hadoop-3.2.3\share\hadoop\common\lib\guava-27.0-jre.jar;C:\hadoop-
e\hadoop\common\lib\hadoop-auth-3.2.3.jar;C:\hadoop-3.2.3\share\hadoop\common\lib\htrace-
```

```
C:\WINDOWS\system32>cd C:\hadoop-3.2.3\sbin

C:\hadoop-3.2.3\sbin>start-all.cmd
This script is Deprecated. Instead use start-dfs.cmd and start-yarn.cmd
starting yarn daemons

C:\hadoop-3.2.3\sbin>
```

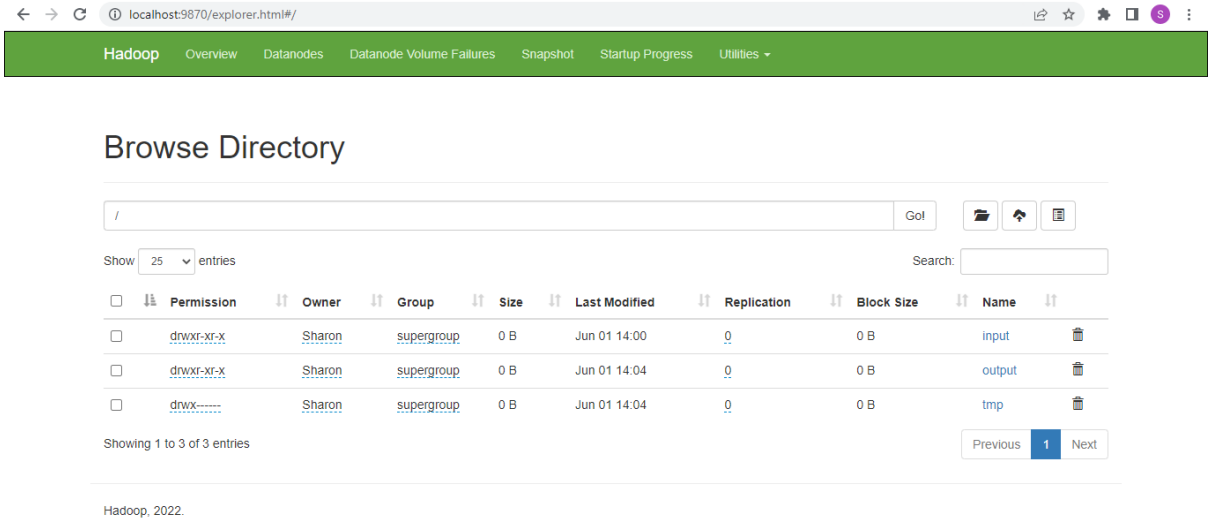

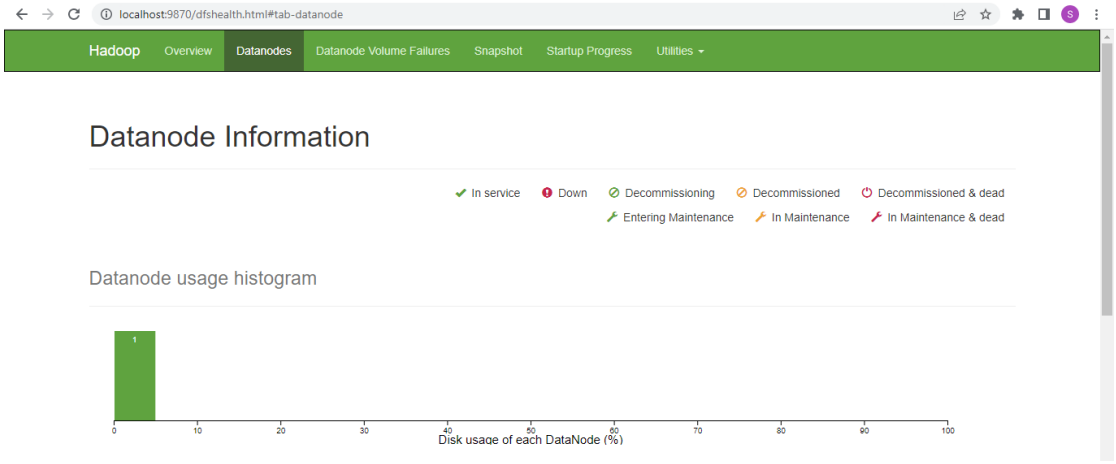


```
C:\WINDOWS\system32>hadoop fs -help
Usage: hadoop fs [generic options]
        [-appendToFile <localsrc> ... <dst>]
        [-cat [-ignoreCrc] <src> ...]
        [-checksum <src> ...]
        [-chgrp [-R] GROUP PATH...]
        [-chmod [-R] <MODE[,MODE]... | OCTALMODE> PATH...]
        [-chown [-R] [OWNER][:[GROUP]] PATH...]
        [-copyFromLocal [-f] [-p] [-l] [-d] [-t <thread count>] [-q <thread pool queue
        [-copyToLocal [-f] [-p] [-crc] [-ignoreCrc] [-t <thread count>] [-q <thread poo
        [-count [-q] [-h] [-v] [-t [<storage type>]] [-u] [-x] [-e] <path> ...]
        [-cp [-f] [-p | -p[topax]] [-d] [-t <thread count>] [-q <thread pool queue size
        [-createSnapshot <snapshotDir> [<snapshotName>]]
        [-deleteSnapshot <snapshotDir> <snapshotName>]
        [-df [-h] [<path> ...]]
        [-du [-s] [-h] [-v] [-x] <path> ...]
        [-expunge [-immediate]]
        [-find <path> ... <expression> ...]
        [-get [-f] [-p] [-crc] [-ignoreCrc] [-t <thread count>] [-q <thread pool queue
        [-getfacl [-R] <path>]
        [-getfattr [-R] {-n name | -d} [-e en] <path>]
        [-getmerge [-nl] [-skip-empty-file] <src> <localdst>]
        [-head <file>]
        [-help [cmd ...]]
        [-ls [-C] [-d] [-h] [-q] [-R] [-t] [-S] [-r] [-u] [-e] [<path> ...]]
        [-mkdir [-p] <path> ...]
```

Exercise 1.2: Basic Hadoop operations

1. Check Hadoop version: hadoop version

```
C:\WINDOWS\system32>hadoop version
Hadoop 3.2.3
Source code repository https://github.com/apache/hadoop -r abe5358143720085498613d399be3bbf01e0f131
Compiled by ubuntu on 2022-03-20T01:18Z
Compiled with protoc 2.5.0
From source with checksum 39bb14faec14b3aa25388a6d7c345fe8
This command was run using /C:/hadoop-3.2.3/share/hadoop/common/hadoop-common-3.2.3.jar

C:\WINDOWS\system32>
```

2. List files in HDFS: hadoop fs -ls /

```
C:\WINDOWS\system32>hadoop fs -ls /
Found 3 items
drwxr-xr-x   - Sharon supergroup          0 2022-06-01 14:00 /input
drwxr-xr-x   - Sharon supergroup          0 2022-06-01 14:04 /output
drwx------   - Sharon supergroup          0 2022-06-01 14:04 /tmp

C:\WINDOWS\system32>
```

3. Create a hadoopdemo directory: hadoop fs -mkdir /hadoopdemo

```
C:\WINDOWS\system32>hadoop fs -mkdir /hadoopdemo

C:\WINDOWS\system32>
```

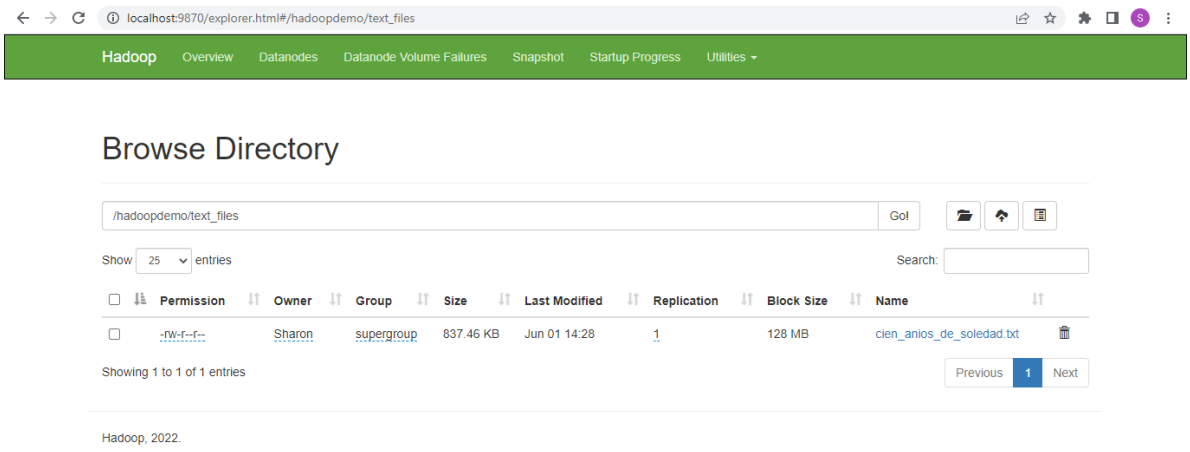

4. Create several sub-directories nested in hadoopdemo, e.g. text files, raw data

7. Remove the sub-directory hadoop fs -rm -r /hadoopdemo/text files

```
C:\WINDOWS\system32>hadoop fs -rm -r /hadoopdemo/text_files
Deleted /hadoopdemo/text_files

C:\WINDOWS\system32>
```

8. Change the content of file.txt in the local system and overwrite it in Hadoop hadoop fs -put -f file.txt /hadoopdemo/text files

9. Read the content of the file: hadoop fs -cat /hadoopdemo/text files/file.txt

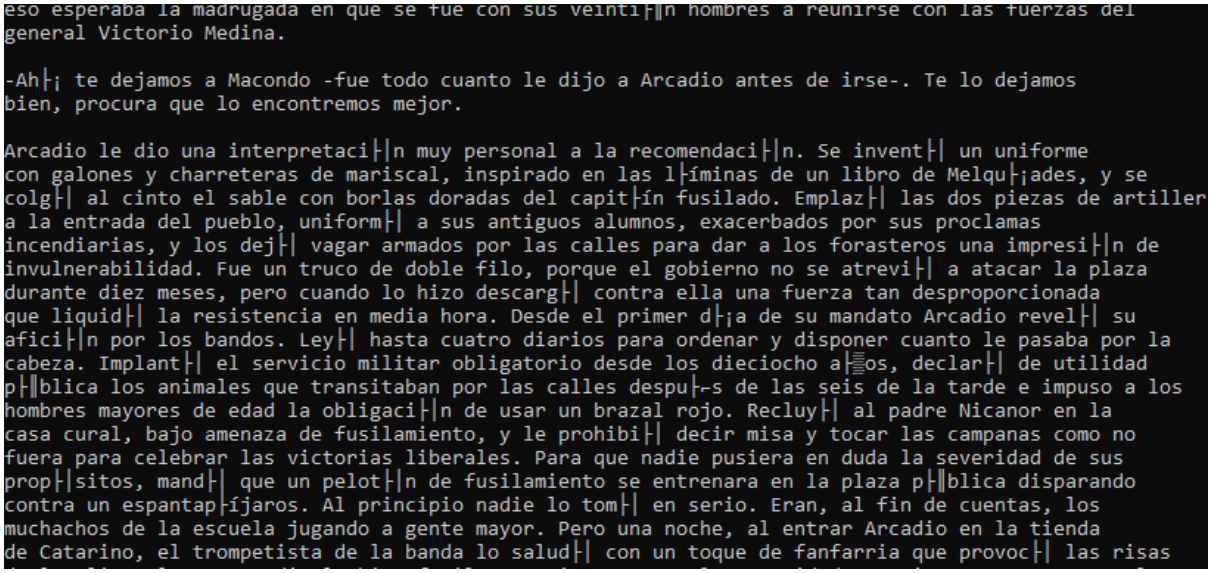

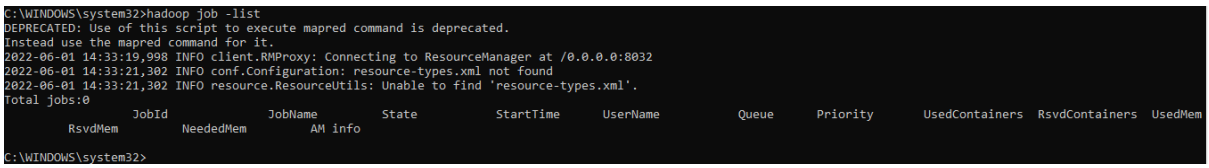

5. Transfer and store a data file from local systems to Hadoop: hadoop fs -put file.txt /hadoopdemo/text files

Exercise 1.3: WordCount MapReduce example

1. First a directory is created named word_count using command: hadoop fs -mkdir word_count
2. Then text dataset used in this exercise is put in the folder using command: hadoop fs -put D:\OneDrive\Desktop\file101.txt /word_count
3. Now we will use the map-reduce that already comes with the Hadoop installation and is present at the path “C:\hadoop-3.2.3\share\hadoop\mapreduce” and help to load the word count program
4. The output of the program is stored in /word_count /word_count_output and from there we can see count of individual words at the URL “http://localhost:9870/explorer.html#/word_count/output1”

Commands Used are:

hadoop fs -put D:\OneDrive\Desktop\file101.txt /word_count

hadoop jar C:\hadoop-3.2.3\share\hadoop\mapreduce\hadoop-mapreduce-examples-3.2.3.jar wordcount /word_count /word_count_output

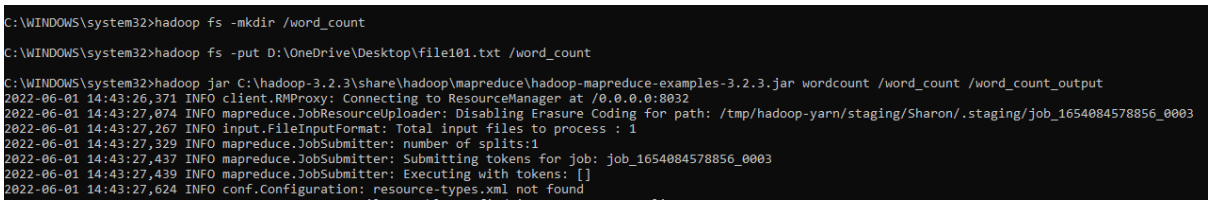

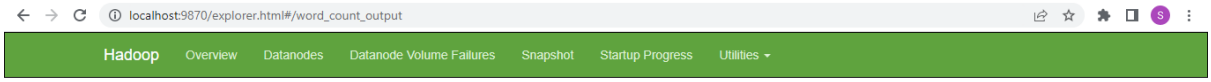

Browse Directory

/word_count_output Go!

Show 25 entries Search:

Permission	Owner	Group	Size	Last Modified	Replication	Block Size	Name
-rw-r--r--	Sharon	supergroup	0 B	Jun 01 14:43	1	128 MB	_SUCCESS
-rw-r--r--	Sharon	supergroup	228.16 KB	Jun 01 14:43	1	128 MB	part-r-00000

Showing 1 to 2 of 2 entries Previous 1 Next

Hadoop, 2022.

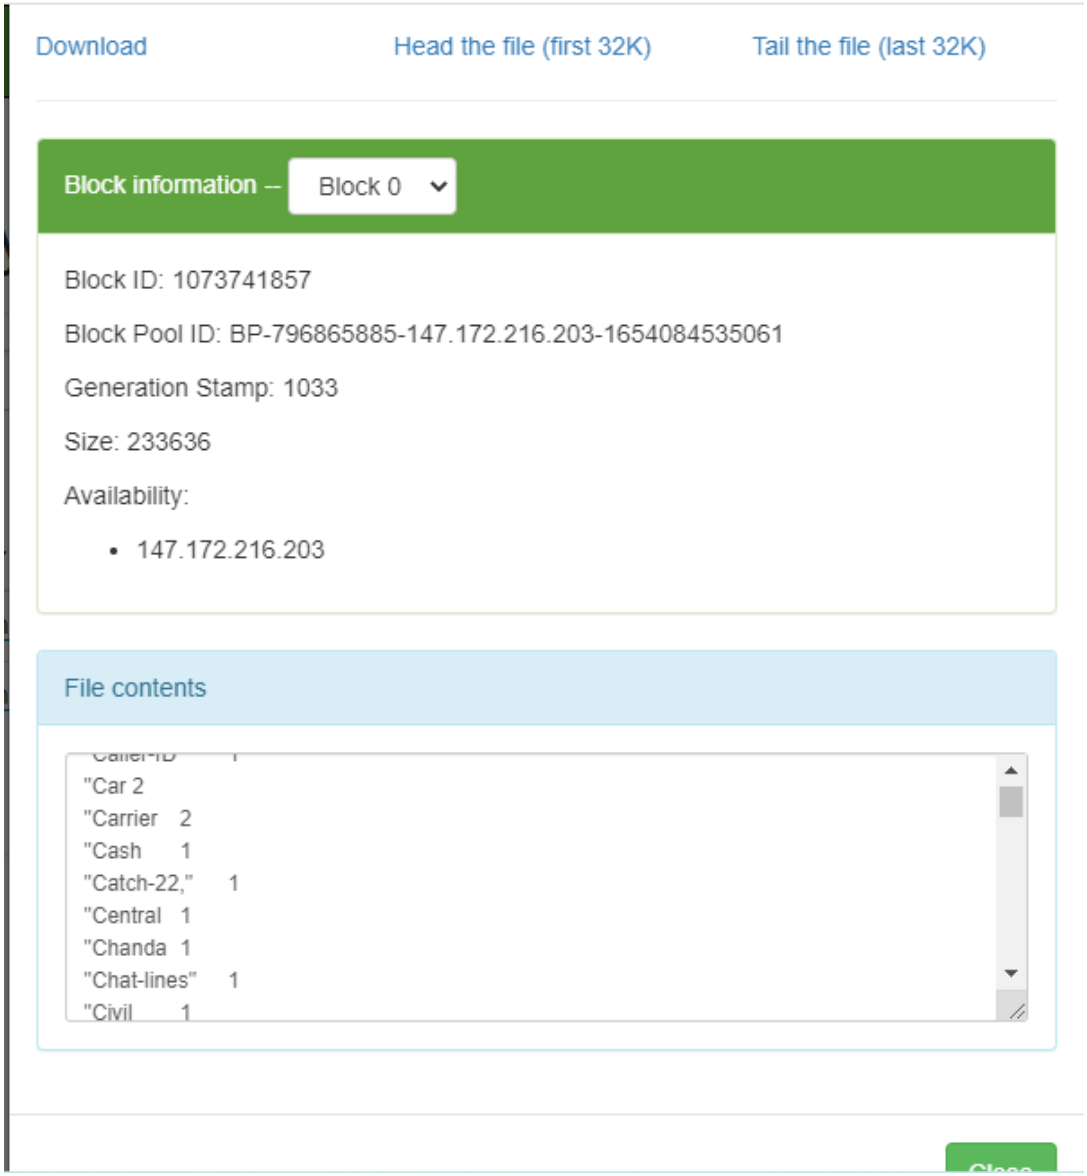

Exercise 2: Analysis of Airport efficiency with Map Reduce

1. First a directory is created named “airport_efficiency” using command: hadoop fs -mkdir /airport_efficiency
2. Then csv file is put in the folder using command: hadoop fs -put D:\OneDrive\Desktop\Data Analytics\DDA LAB\Lab 6\T_ONTIME_REPORTING.csv /airport_efficiency
3. Now we will use custom mapper and reducer functions
4. The output of the program is stored in /airport_efficiency/output2 and from there we can see output at the URL “http://localhost:9870/explorer.html#/airport_efficiency/output”

```
C:\WINDOWS\system32>hadoop fs -mkdir /airport_efficiency
```

```
C:\WINDOWS\system32>hadoop fs -put "D:\OneDrive\Desktop\Data Analytics\DDA LAB\Lab 6\T_ONTIME_REPORTING.csv" /airport_efficiency
```

1. Computing the maximum, minimum, and average departure delay for each airport

- For the mapper function each line is read from the csv file and split into tokens.
- Attributes Origin and Departure Delay is read and printed to the reducer.
- The reducer takes this input and checks whether departure delay has a value.
- Dictionary dep_delay stores origin as the key and delay as the value. This ensures unique airports.
- The key of the dictionary dep_delay is iterated and average, min and max values are found and printed on the final output console.

```
#Reducer
for dep in dep_delay.keys():
    ave_dep = sum(dep_delay[dep])*1.0 / len(dep_delay[dep])
    min_dep = min(dep_delay[dep])
    max_dep = max(dep_delay[dep])
    print ('%s\t%s\t%s\t%s'% (dep, ave_dep,min_dep,max_dep))
```

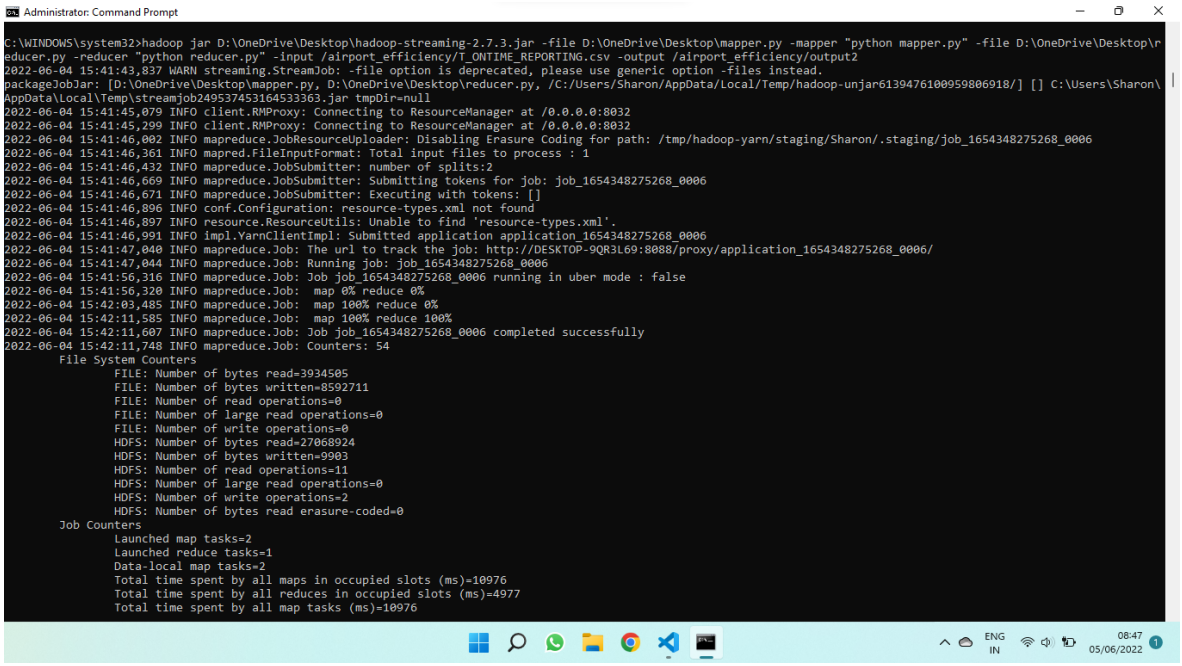

Airport, Average, Minimum, Maximum

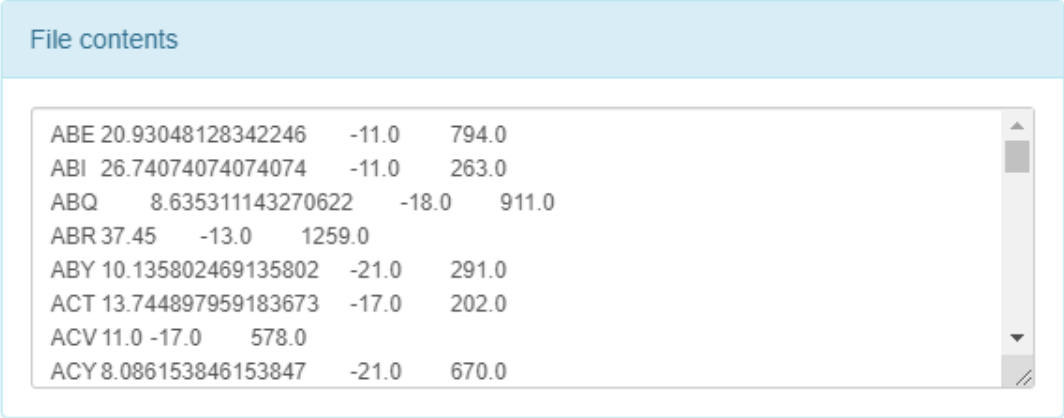

2. Computing a ranking list that contains top 10 airports by their average Arrival delay.

- For the mapper function each line is read from the csv file and split into tokens.
- Attributes Destination and Arrival Delay is read and printed to the reducer.
- The reducer takes this input and checks whether Arrival delay column has a value.
- Dictionary arrival_delay stores destination as the key and delay as the value. This ensures unique airports.
- The key of the dictionary arrival_delay is iterated and average values are found.
- Finally, the average values are sorted in reverse order and the top 10 airports are printed along with their ranks.

```
for dep in arrival_delay.keys():
    ave_dep = sum(arrival_delay[dep])*1.0 / len(arrival_delay[dep])
    avg_delay.append((dep,ave_dep))

final_list=[]
final_list=sorted(avg_delay, key=itemgetter(1),reverse=True)
#Reducer
for rank,delay in enumerate(final_list):
    if rank<10:
        print ('%s\t%s\t%s'% (rank+1,delay[0],delay[1]))
```

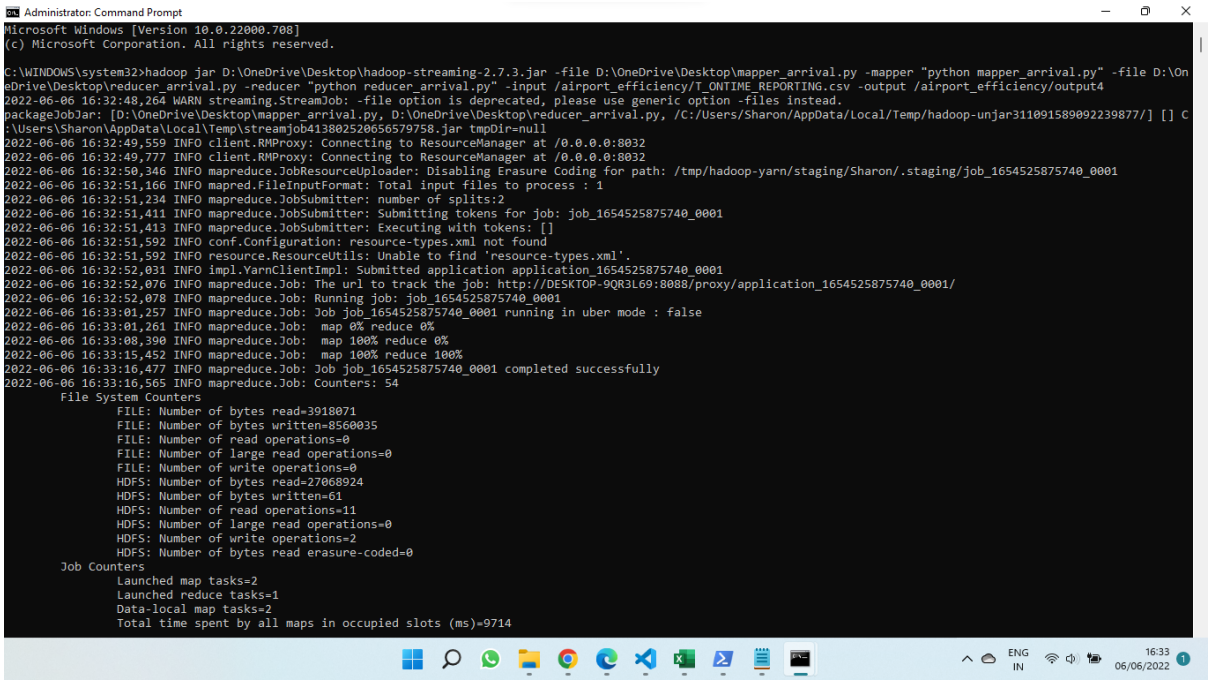

Output: Ranking of Airports

```
1    ELM 88.92307692307692
2    GGG 61.5
3    BPT 59.6
4    BMI 39.825688073394495
5    LWS 39.38775510204081
6    LAW 37.8
7    GRB 33.23131672597865
8    ESC 32.604651162790695
9    ABI 29.296296296296298
10   OAJ 28.654761904761905
```

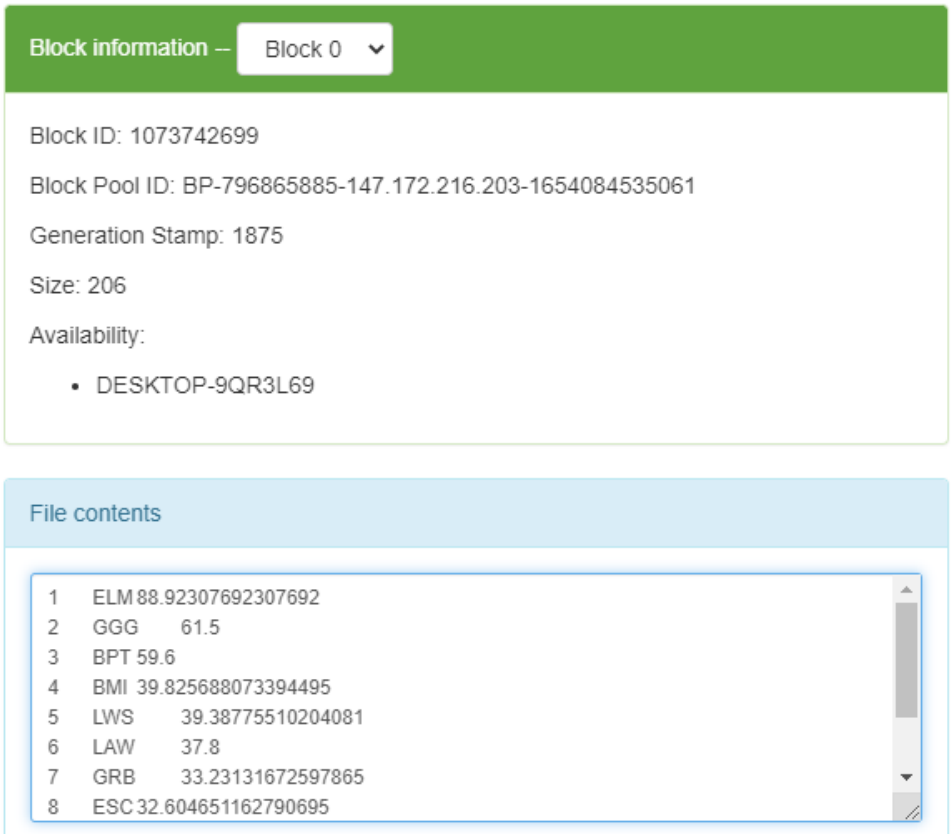

Exercise 3: Analysis of Movie dataset using Map and Reduce

ratings.dat - UserID::MovieID::Rating::Timestamp

movies.dat - MovieID::Title::Genres

1. Find the movie title which has the maximum average rating?

- For the mapper function each line is read from the ratings.dat file and split into user_id,movie_id,rating,timestamp. Movie_Id and rating is printed to the reducer.
- The reducer takes this input and checks whether each movie has a rating
- Dictionary movie_list stores movie as the key and rating as the value. This ensures unique movies.
- The dictionary movie_list is iterated based on key: “movies” and average values are found.
- Finally, the average values are sorted in reverse order and movie with max average rating is found out.
- The title of the movie is found from the “movies.dat” file and movies with max avg rating is printed.

```python
for movie in movie_list.keys():
    avg_rate= sum(movie_list[movie])*1.0 / len(movie_list[movie])
    avg_rating.append((movie,avg_rate))

final_list=[]
final_list=sorted(avg_rating, key=itemgetter(1),reverse=True)
max_value=max(final_list,key=itemgetter(1))
#Reducer
for movie_rating in final_list:
    if movie_rating[1]==max_value[1]:
        print(f"{df.loc[df.movie_id==int(movie_rating[0])]['title'].values[0]}\t{movie_rating[1]}")
```

Commands Used:

hadoop fs -put D:\OneDrive\Desktop\Data Analytics\DDA LAB\Lab 6\ml-10m\ml-10M100K\ratings.dat /movie_data

hadoop jar D:\OneDrive\Desktop\hadoop-streaming-2.7.3.jar -file D:\OneDrive\Desktop\movie_data\avg_rating_mapper.py -mapper "python avg_rating_mapper.py" -file D:\OneDrive\Desktop\movie_data\avg_rating_reducer.py -reducer "python avg_rating_reducer.py" -input /movie_data/ratings.dat -output /movie_data/output1

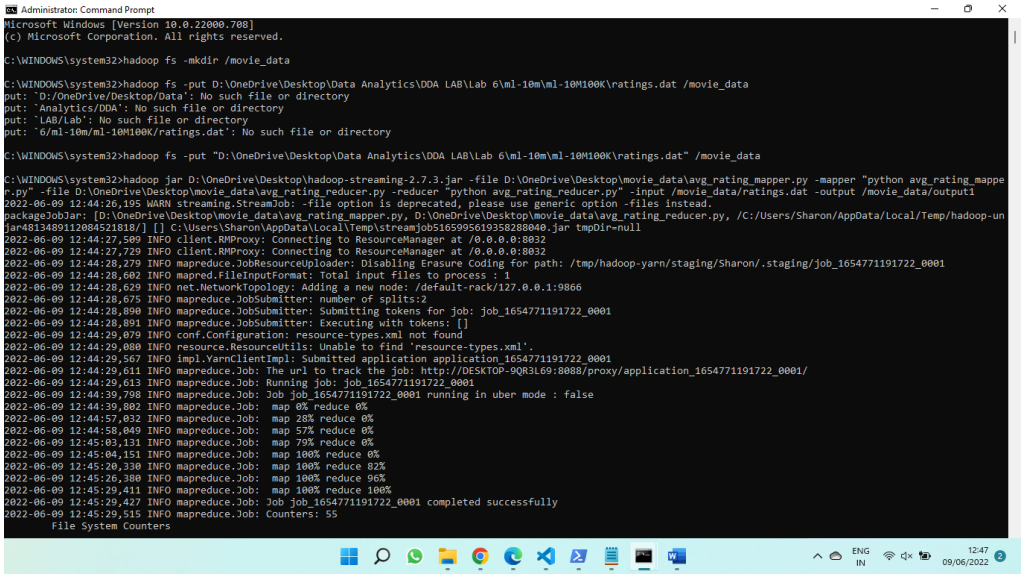

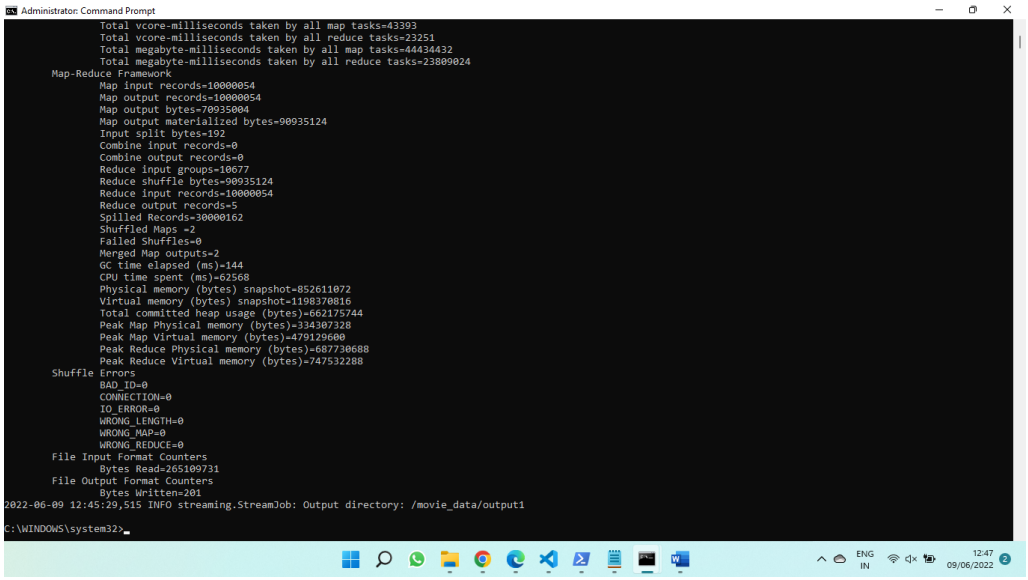

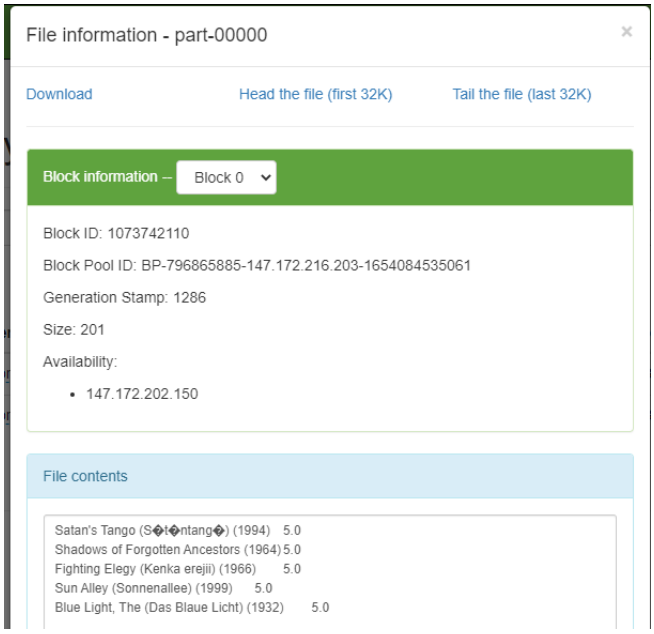

2. Find the user who has assign lowest average rating among all the users who rated more than 40 times?

- For the mapper function each line is read from the ratings.dat file and split into user_id,movie_id,rating,timestamp. user_id and rating is printed to the reducer.
- The reducer takes this input and checks whether each user has a rating
- Dictionary user_list stores users as the key and rating as the value. This ensures unique users.
- The dictionary user_list is iterated based on key: “user” and users who have rated more than 40 times are found.
- Then, the average values are found from user_list1 and min value is computed.
- Finally, user who has assign lowest average rating is printed as the output

```
for user in user_list.keys():
    if len(user_list[user])>40:
        user_list1[user]=user_list[user]

for user in user_list1.keys():
    avg_rating[user]= sum(user_list1[user])*1.0 / len(user_list1[user])

min_value=min(avg_rating.values())
#Reducer
for user in user_list1.keys():
    if avg_rating[user]==min_value:
        print ('%s\t%s'% (user,avg_rating[user]))
```

Commands Used:

hadoop jar D:\OneDrive\Desktop\hadoop-streaming-2.7.3.jar -file D:\OneDrive\Desktop\movie_data\avg_rating_user_mapper.py -mapper "python avg_rating_user_mapper.py" -file D:\OneDrive\Desktop\movie_data\avg_rating_user_reducer.py -reducer "python avg_rating_user_reducer.py" -input /movie_data/ratings.dat -output /movie_data/output2

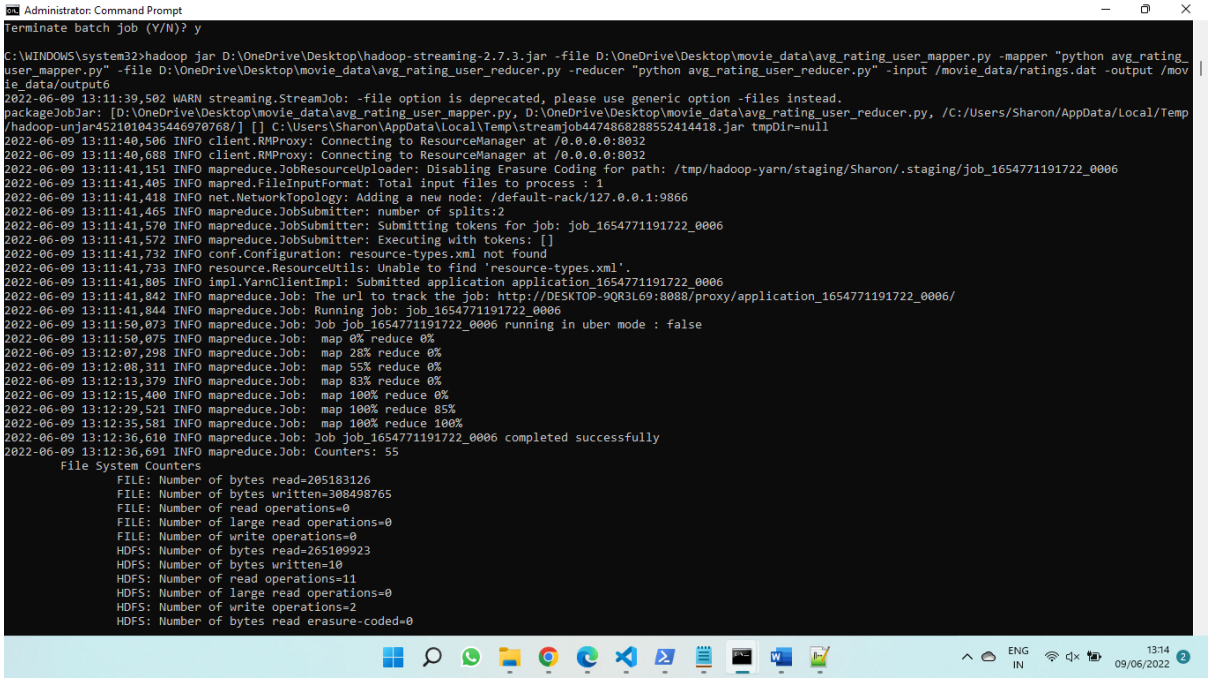

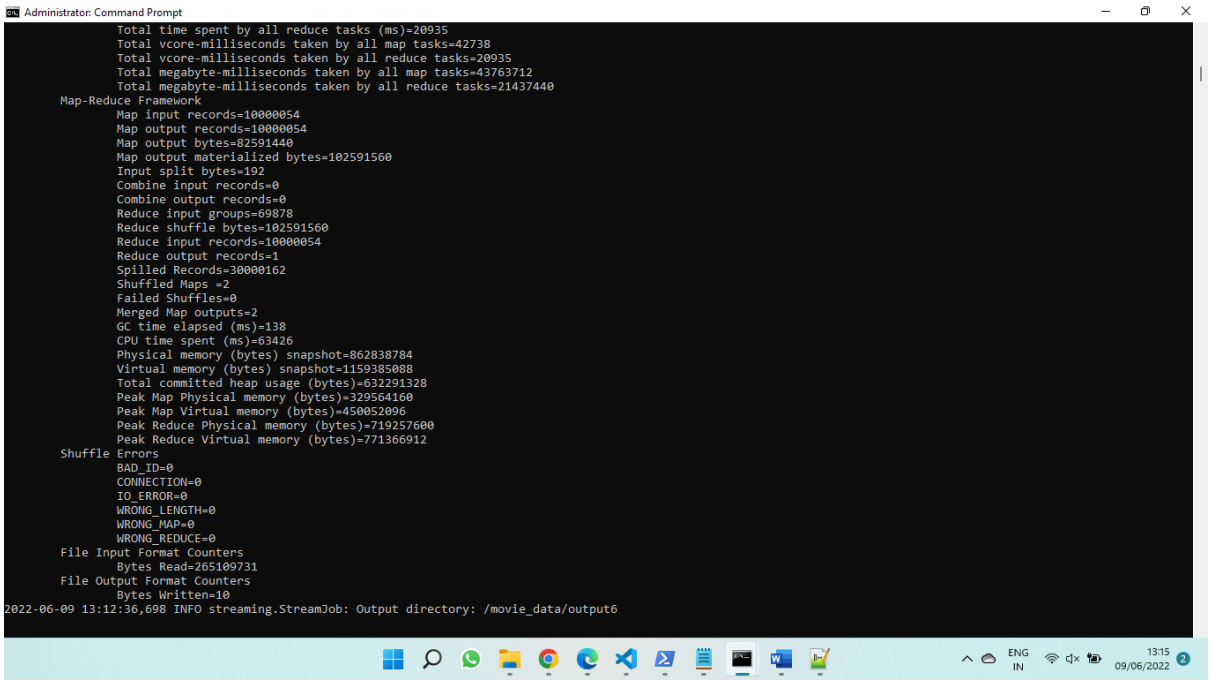

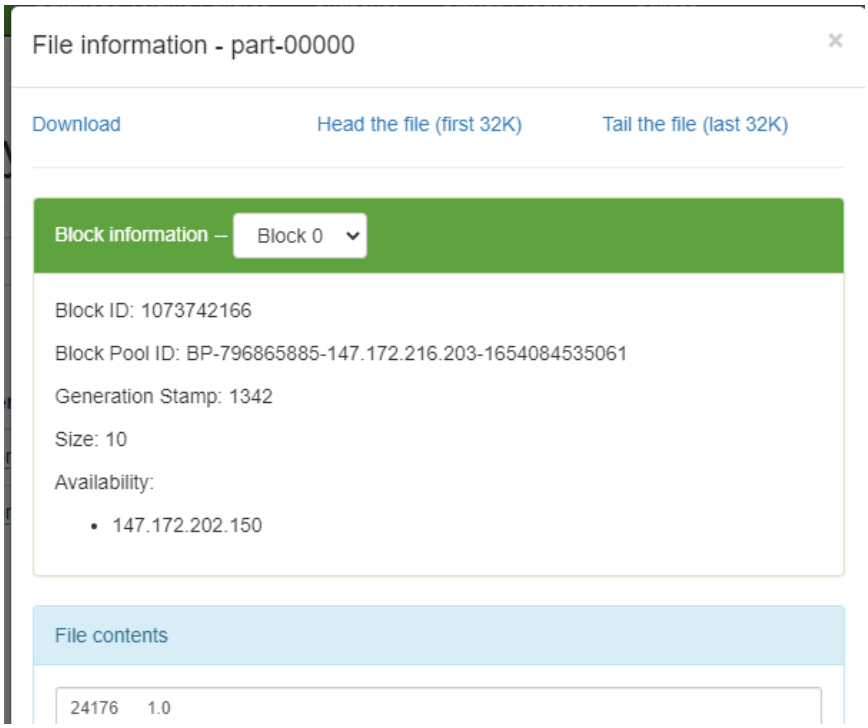

User 24176 assigns the lowest average rating of 1 among all the users.

3. Find the highest average rated movie genre?

- As highest average rated genre is to be found out, preprocessing step needs to be carried out.
- In the preprocessing step, movies.dat and ratings.dat is combined and attributes “genre”,”rating” is selected from the dataframe. This is written to genre.csv file

- For the mapper function each line is read from the “genre.csv” file and split into genre,rating which is printed to the reducer.
- The reducer takes this input and checks whether each genre has a rating
- Dictionary genre_list stores genre as the key and rating as the value. This ensures unique genres.
- The dictionary genre_list is iterated based on key: “genre” and average values are found.
- Then, from the average values the maximum value is computed.
- Finally, genre which has highest average rating is given as the output

```
for genre in genre_list.keys():
    avg_rating[genre]= sum(genre_list[genre])*1.0 / len(genre_list[genre])

max_value=max(avg_rating.values())

#Reducer
for genre in genre_list.keys():
    if avg_rating[genre]==max_value:
        print ('%s\t%s'% (genre,max_value))
```

Commands Used:

hadoop fs -put "D:\OneDrive\Desktop\Data Analytics\DDA LAB\Lab 6\genre_rate1.csv" /movie_data

hadoop jar D:\OneDrive\Desktop\hadoop-streaming-2.7.3.jar -file D:\OneDrive\Desktop\movie_data\genre_mapper.py -mapper "python genre_mapper.py" -file D:\OneDrive\Desktop\movie_data\genre_reducer.py -reducer "python genre_reducer.py" -input /movie_data/genre_rate1.csv -output /movie_data/output15

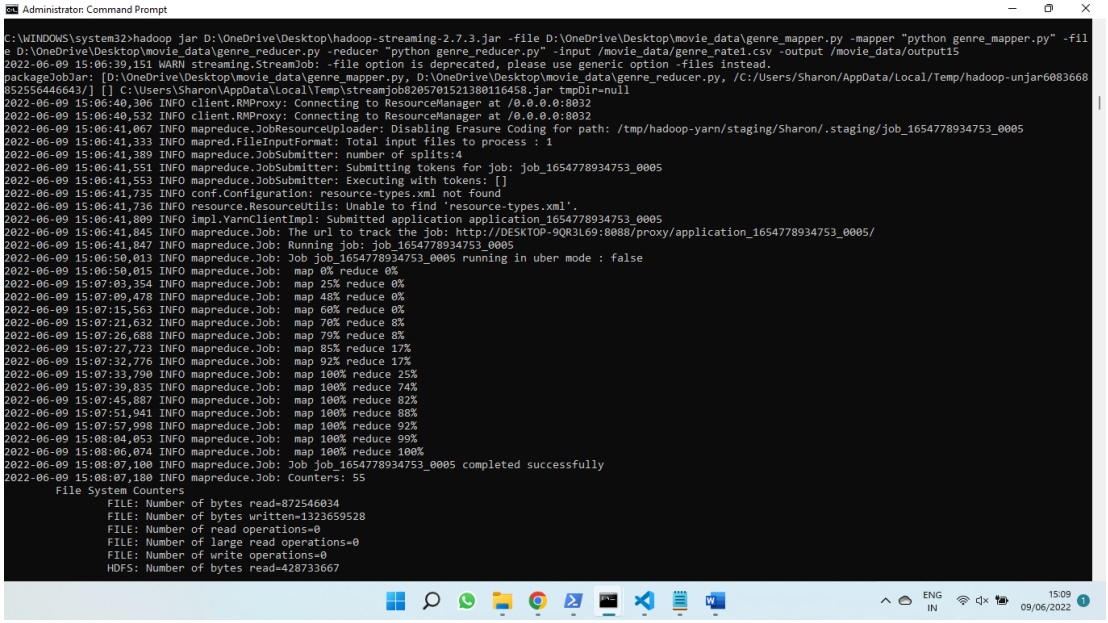

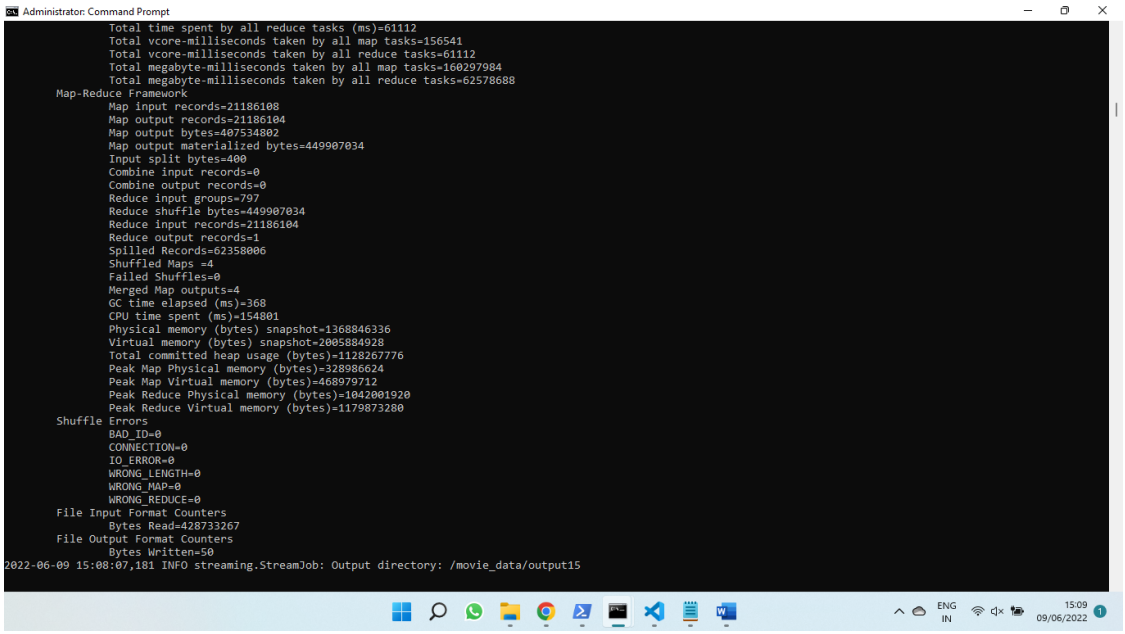

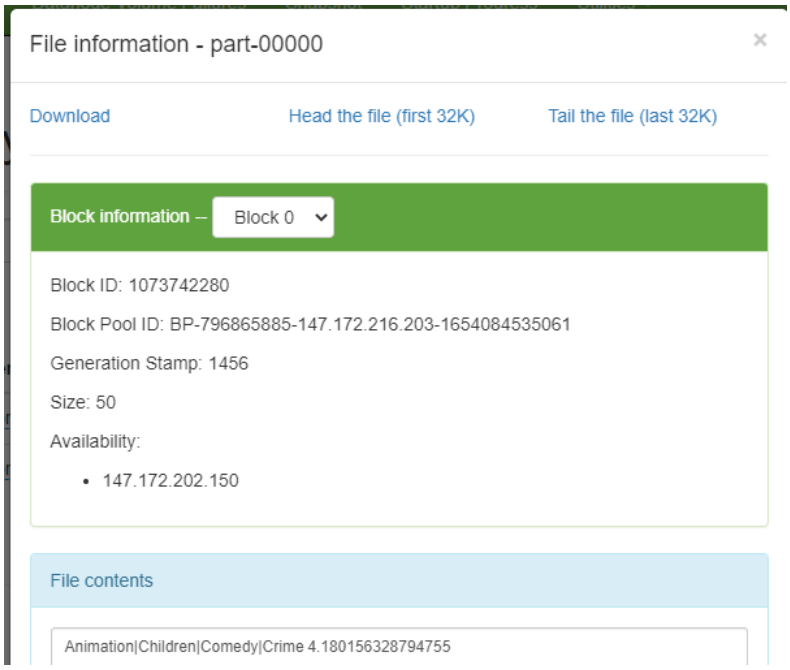

Task a: Performance Analysis

Mapper * Reducer =2*1=CPU time spent (ms)=67521
Mapper * Reducer =4*1= CPU time spent (ms)=45880
Mapper * Reducer =8*1= CPU time spent (ms)=53266
Mapper * Reducer =16*1=CPU time spent (ms)=72717

No of Mapper * Reducer	2*1	4*1	8*1	16*1
Time (sec)	67.521	45.880	53.266	72.717

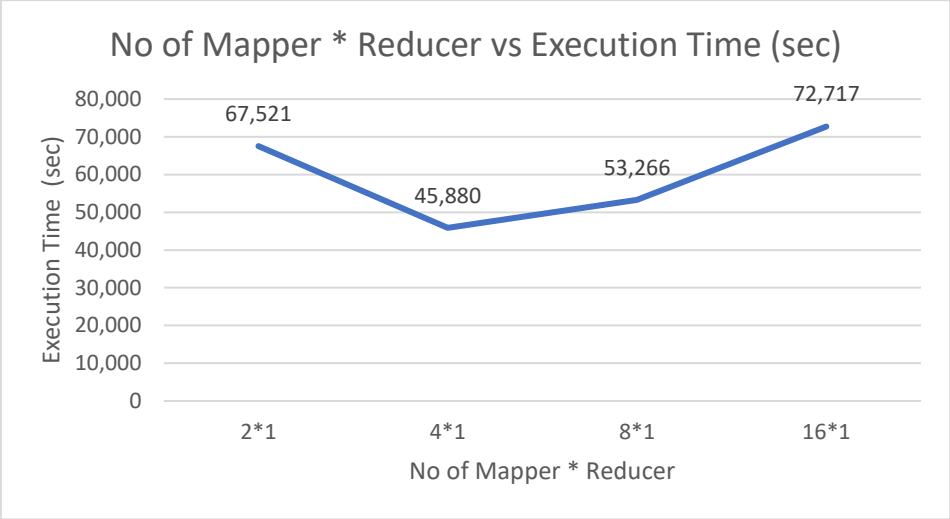

Performance Analysis: When the number of mapper * reducer is varied from 2*1 to 4*1, this tells us that as mappers are increased, number of splits in data gets increased. This leads to reduction of Execution time. It reduces from 67 sec to 45 sec. But as the number is increased more from 8 *1 to 16*1 , the time increases. This tells making more number of splits actually creates performance overhead, which lead to degradation of performance. Here time is indicated in sec and graph shows german numbering format

Task b: Performance Analysis

Mapper * Reducer =2*1=CPU time spent (ms)=64491
Mapper * Reducer =4*1=CPU time spent (ms)=45860
Mapper * Reducer =8*1=CPU time spent (ms)=61073
Mapper * Reducer =16*1=CPU time spent (ms)=71733

No of Mapper * Reducer	2*1	4*1	8*1	16*1
Time (sec)	64.491	45.860	61.073	71.733

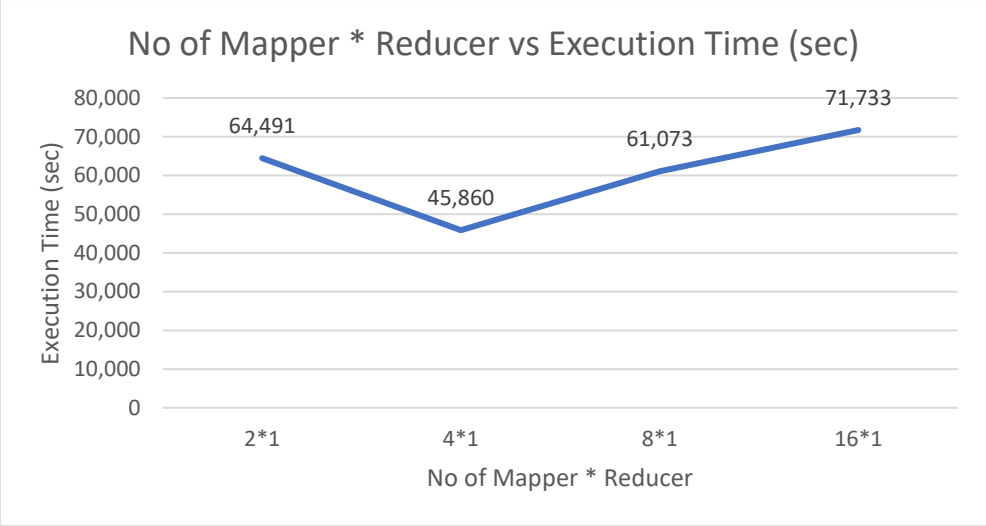

Performance Analysis: When the number of mapper * reducer is varied from 2*1 to 4*1, this tells us that as mappers are increased, number of splits in data gets increased. This leads to reduction of Execution time. But as the number is increased more from 8 *1 to 16*1 , the time increases. This tells making more number of splits actually creates performance overhead, which lead to degradation of performance.

Task c: Performance Analysis

Mapper * Reducer =2*1=CPU time spent (ms)=173176
Mapper * Reducer =4*1= CPU time spent (ms)=167630
Mapper * Reducer =8*1=CPU time spent (ms)=174853
Mapper * Reducer =16*1=CPU time spent (ms)=108844

No of Mapper * Reducer	2*1	4*1	8*1	16*1
Time (sec)	173.176	167.630	174.853	108.844

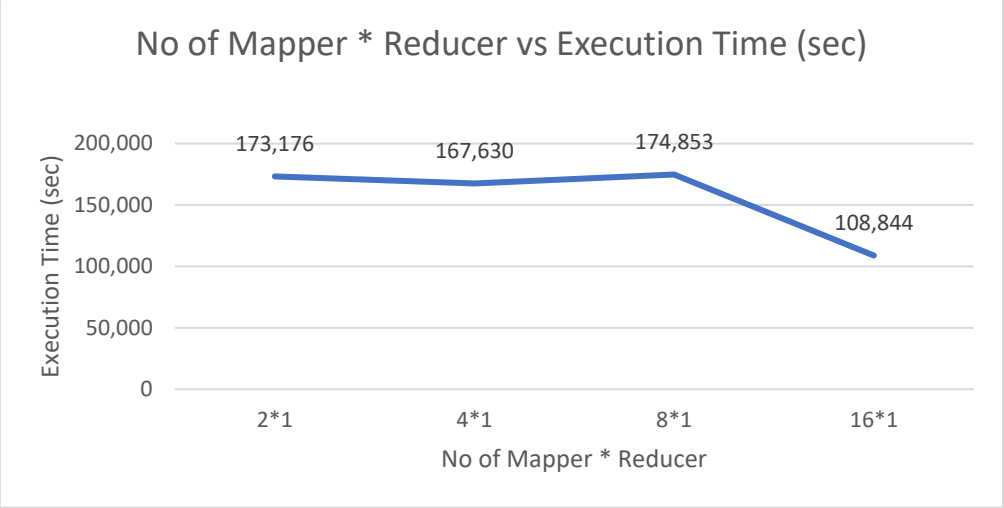

Performance Analysis: When the number of mapper * reducer is varied from 2*1 to 4*1, this tells us that as mappers are increased, number of splits in data gets increased. This leads to reduction of Execution time. As the number is increased more from 8 *1 to 16*1 , the time further decreases. As the dataset size is huge, having 16 mappers is advantageous. The performance increases with 16 mappers by a huge margin. Having a distributed architecture leads to speedup. Least time is with 16 mappers and 1 reducer which is 108 sec.

References:

https://www.geeksforgeeks.org/hadoop-streaming-using-python-word-count-problem/

https://stackoverflow.com/questions/47618867/how-to-run-python-mapreduce-in-hadoop-streaming

https://medium.com/@pedro.a.hdez.a/hadoop-3-2-2-installation-guide-for-windows-10-454f5b5c22d3#d3ca